

BA-PROJECT_GROUP-8

2023-04-25

Installing and Loading the required packages

```
library(ISLR)
library(caret)

## Loading required package: ggplot2
## Loading required package: lattice
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(class)
library(e1071)
library(tidyverse)

## — Attaching core tidyverse packages ————— tidyverse
2.0.0 —
## ✓ forcats   1.0.0   ✓ stringr   1.5.0
## ✓ lubridate 1.9.2   ✓ tibble    3.2.1
## ✓ purrr     1.0.1   ✓ tidyr     1.3.0
## ✓ readr     2.1.4

## — Conflicts —————
tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
## ✗ purrr::lift()    masks caret::lift()
## ⓘ Use the conflicted package (<http://conflicted.r-lib.org/>) to force all
conflicts to become errors

library(ggplot2)
library(ggcorrplot)
library(rattle)
```

```
## Loading required package: bitops
## Rattle: A free graphical interface for data science with R.
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.

library(gmodels)
library(modelr)
library(Hmisc)

##
## Attaching package: 'Hmisc'
##
## The following object is masked from 'package:e1071':
##
##     impute
##
## The following objects are masked from 'package:dplyr':
##
##     src, summarize
##
## The following objects are masked from 'package:base':
##
##     format.pval, units

library(missForest)
library(pROC)

## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
##
## The following object is masked from 'package:gmodels':
##
##     ci
##
## The following objects are masked from 'package:stats':
##
##     cov, smooth, var

library(ROCR)
library(cutpointr)

##
## Attaching package: 'cutpointr'
##
## The following objects are masked from 'package:pROC':
##
##     auc, roc
##
## The following objects are masked from 'package:caret':
```

```
##
##      precision, recall, sensitivity, specificity

library(ROSE)

## Loaded ROSE 0.0-4

library(cowplot)

##
## Attaching package: 'cowplot'
##
## The following object is masked from 'package:lubridate':
##
##      stamp

library(tinytex)
library(caTools)
library(broom)

##
## Attaching package: 'broom'
##
## The following object is masked from 'package:modelr':
##
##      bootstrap

library(rpart)
library(rpart.plot)
```

Loading the Churn dataset to R environment

```
#Included both csv file and .Rdata file in the same working directory where
Rmd file is saved
Raw_Churn_data<- read.csv("Churn_Train.csv")

load(file = "Customers_To_Predict.RData",envir = globalenv())
```

Attributes in the dataset

```
colnames(Raw_Churn_data)

## [1] "state" "account_length"
## [3] "area_code" "international_plan"
## [5] "voice_mail_plan" "number_vmail_messages"
## [7] "total_day_minutes" "total_day_calls"
## [9] "total_day_charge" "total_eve_minutes"
## [11] "total_eve_calls" "total_eve_charge"
## [13] "total_night_minutes" "total_night_calls"
## [15] "total_night_charge" "total_intl_minutes"
## [17] "total_intl_calls" "total_intl_charge"
## [19] "number_customer_service_calls" "churn"
```

```
str(Raw_Churn_data)

## 'data.frame':    3333 obs. of  20 variables:
## $ state                : chr  "NV" "HI" "DC" "HI" ...
## $ account_length       : int   125 108 82 NA 83 89 135 28 86 65
## ...
## $ area_code            : chr   "area_code_510" "area_code_415"
## "area_code_415" "area_code_408" ...
## $ international_plan   : chr   "no" "no" "no" "no" ...
## $ voice_mail_plan      : chr   "no" "no" "no" "yes" ...
## $ number_vmail_messages : int    0 0 0 30 0 0 0 0 0 ...
## $ total_day_minutes     : num   2013 292 300 110 337 ...
## $ total_day_calls       : int    99 99 109 71 120 81 81 87 115 137
## ...
## $ total_day_charge      : num    28.7 49.6 51 18.8 57.4 ...
## $ total_eve_minutes     : num   1108 221 181 182 227 ...
## $ total_eve_calls       : int    107 93 100 108 116 74 114 92 112 83
## ...
## $ total_eve_charge      : num    14.9 18.8 15.4 15.5 19.3 ...
## $ total_night_minutes   : num    243 229 270 184 154 ...
## $ total_night_calls     : int     92 110 73 88 114 120 82 112 95 111
## ...
## $ total_night_charge    : num    10.95 10.31 12.15 8.27 6.93 ...
## $ total_intl_minutes    : num     10.9 14 11.7 11 15.8 9.1 10.3 10.1
## 9.8 12.7 ...
## $ total_intl_calls      : int      7 9 4 8 7 4 6 3 7 6 ...
## $ total_intl_charge     : num     2.94 3.78 3.16 2.97 4.27 2.46 2.78
## 2.73 2.65 3.43 ...
## $ number_customer_service_calls: int     0 2 0 2 0 1 1 3 2 4 ...
## $ churn                 : chr   "no" "yes" "yes" "no" ...
```

Removing the columns that are not needed

```
Churn_Train_Data<- Raw_Churn_data[, -c(1:3)]
colnames(Churn_Train_Data)

## [1] "international_plan"      "voice_mail_plan"
## [3] "number_vmail_messages"  "total_day_minutes"
## [5] "total_day_calls"        "total_day_charge"
## [7] "total_eve_minutes"      "total_eve_calls"
## [9] "total_eve_charge"       "total_night_minutes"
## [11] "total_night_calls"      "total_night_charge"
## [13] "total_intl_minutes"     "total_intl_calls"
## [15] "total_intl_charge"      "number_customer_service_calls"
## [17] "churn"
```

Converting international plan, voice mail plan and churn variables to binary

```
Churn_Train_Data$international_plan<-
ifelse(Churn_Train_Data$international_plan == "yes",1,0)
```

```
Churn_Train_Data$voice_mail_plan<- ifelse(Churn_Train_Data$voice_mail_plan
=="yes",1,0)
```

```
Churn_Train_Data$churn<- ifelse(Churn_Train_Data$churn == "yes",1,0)
```

Verify any NA values present in the dataset

```
any(is.na.data.frame(Churn_Train_Data))
```

```
## [1] TRUE
```

```
colMeans(is.na(Churn_Train_Data))*100
```

```
##          international_plan          voice_mail_plan
##          0.000000          0.000000
##    number_vmail_messages    total_day_minutes
##          6.000600          6.000600
##          total_day_calls    total_day_charge
##          6.000600          6.000600
##          total_eve_minutes    total_eve_calls
##          9.030903          6.000600
##          total_eve_charge    total_night_minutes
##          6.000600          6.000600
##          total_night_calls    total_night_charge
##          0.000000          6.000600
##          total_intl_minutes    total_intl_calls
##          6.000600          9.030903
##          total_intl_charge number_customer_service_calls
##          6.000600          6.000600
##          churn
##          0.000000
```

Review All the columns after data manipulation and attribute exclusion

```
summary(Churn_Train_Data)
```

```
## international_plan voice_mail_plan number_vmail_messages
total_day_minutes
## Min.   :0.00000   Min.   :0.0000   Min.   : -10.000   Min.   :  0.0
## 1st Qu.:0.00000   1st Qu.:0.0000   1st Qu.:  0.000   1st Qu.: 149.3
## Median :0.00000   Median :0.0000   Median :  0.000   Median : 190.5
## Mean   :0.09691   Mean   :0.2766   Mean   :  7.333   Mean   : 418.9
## 3rd Qu.:0.00000   3rd Qu.:1.0000   3rd Qu.: 16.000   3rd Qu.: 237.8
## Max.   :1.00000   Max.   :1.0000   Max.   : 51.000   Max.   :2185.1
##                                     NA's   :200       NA's   :200
## total_day_calls total_day_charge total_eve_minutes total_eve_calls
## Min.   :  0.0   Min.   : 0.00   Min.   :  0.0   Min.   :  0.0
## 1st Qu.: 87.0   1st Qu.:24.45   1st Qu.: 170.5   1st Qu.: 87.0
## Median :101.0   Median :30.65   Median : 209.9   Median :100.0
## Mean   :100.3   Mean   :30.63   Mean   : 324.3   Mean   :100.1
## 3rd Qu.:114.0   3rd Qu.:36.84   3rd Qu.: 257.6   3rd Qu.:114.0
```

```
## Max. :165.0 Max. :59.64 Max. :1244.2 Max. :170.0
## NA's :200 NA's :200 NA's :301 NA's :200
## total_eve_charge total_night_minutes total_night_calls total_night_charge
## Min. : 0.00 Min. : 23.2 Min. : 33.0 Min. : 1.040
## 1st Qu.:14.14 1st Qu.:167.3 1st Qu.: 87.0 1st Qu.: 7.530
## Median :17.09 Median :201.4 Median :100.0 Median : 9.060
## Mean :17.08 Mean :201.2 Mean :100.1 Mean : 9.054
## 3rd Qu.:20.00 3rd Qu.:235.3 3rd Qu.:113.0 3rd Qu.:10.590
## Max. :30.91 Max. :395.0 Max. :175.0 Max. :17.770
## NA's :200 NA's :200 NA's :200
## total_intl_minutes total_intl_calls total_intl_charge
## Min. : 0.00 Min. : 0.00 Min. :0.000
## 1st Qu.: 8.50 1st Qu.: 3.00 1st Qu.:2.300
## Median :10.30 Median : 4.00 Median :2.780
## Mean :10.23 Mean : 4.47 Mean :2.762
## 3rd Qu.:12.10 3rd Qu.: 6.00 3rd Qu.:3.270
## Max. :20.00 Max. :20.00 Max. :5.400
## NA's :200 NA's :301 NA's :200
## number_customer_service_calls churn
## Min. :0.000 Min. :0.0000
## 1st Qu.:1.000 1st Qu.:0.0000
## Median :1.000 Median :0.0000
## Mean :1.561 Mean :0.1449
## 3rd Qu.:2.000 3rd Qu.:0.0000
## Max. :9.000 Max. :1.0000
## NA's :200
```

Imputing the missing values with the medians of the columns as the mean value may be very sensitive to outliers.

#Treating the null values with median of the column.

```
Median_ofColumns<- apply(Churn_Train_Data,2,median, na.rm=T)
for (i in colnames(Churn_Train_Data))
Churn_Train_Data[,i][is.na(Churn_Train_Data[,i])]<- Median_ofColumns[i]

any(is.na.data.frame(Churn_Train_Data))#checking for any null values present
after imputation.
```

```
## [1] FALSE
```

```
str(Churn_Train_Data)
```

```
## 'data.frame': 3333 obs. of 17 variables:
## $ international_plan : num 0 0 0 0 0 0 0 0 0 0 ...
## $ voice_mail_plan : num 0 0 0 1 0 0 0 0 0 0 ...
## $ number_vmail_messages : num 0 0 0 30 0 0 0 0 0 0 ...
## $ total_day_minutes : num 2013 292 300 110 337 ...
## $ total_day_calls : num 99 99 109 71 120 81 81 87 115 137
## ...
## $ total_day_charge : num 28.7 49.6 51 18.8 57.4 ...
```

```
## $ total_eve_minutes      : num  1108 221 181 182 227 ...
## $ total_eve_calls        : num   107  93 100 108 116 74 114 92 112 83
...
## $ total_eve_charge       : num   14.9 18.8 15.4 15.5 19.3 ...
## $ total_night_minutes   : num   243 229 270 184 154 ...
## $ total_night_calls      : num    92 110 73 88 114 120 82 112 95 111
...
## $ total_night_charge     : num   10.95 10.31 12.15 8.27 6.93 ...
## $ total_intl_minutes     : num    10.9 14 11.7 11 15.8 9.1 10.3 10.1
9.8 12.7 ...
## $ total_intl_calls       : num    7 9 4 8 7 4 6 3 7 6 ...
## $ total_intl_charge      : num    2.94 3.78 3.16 2.97 4.27 2.46 2.78
2.73 2.65 3.43 ...
## $ number_customer_service_calls: num    0 2 0 2 0 1 1 3 2 4 ...
## $ churn                  : num    0 1 1 0 1 0 0 0 0 1 ...
```

Treating the churn column as numbers can cause issues when using the column in certain functions or models that expect a factor variable. so we are converting the number to factor.

```
Churn_Train_Data$churn<- as.factor(Churn_Train_Data$churn)#converting the integers to factors
```

```
Churn_Train_Data$international_plan<-
as.factor(Churn_Train_Data$international_plan)
```

```
Churn_Train_Data$voice_mail_plan<-
as.factor(Churn_Train_Data$voice_mail_plan)
```

```
#Churn_Train_Data$churn<-
#factor(Churn_Train_Data$churn, levels(Churn_Train_Data$churn)[c(2,1)])
#Changing the order of the churn factor levels.
```

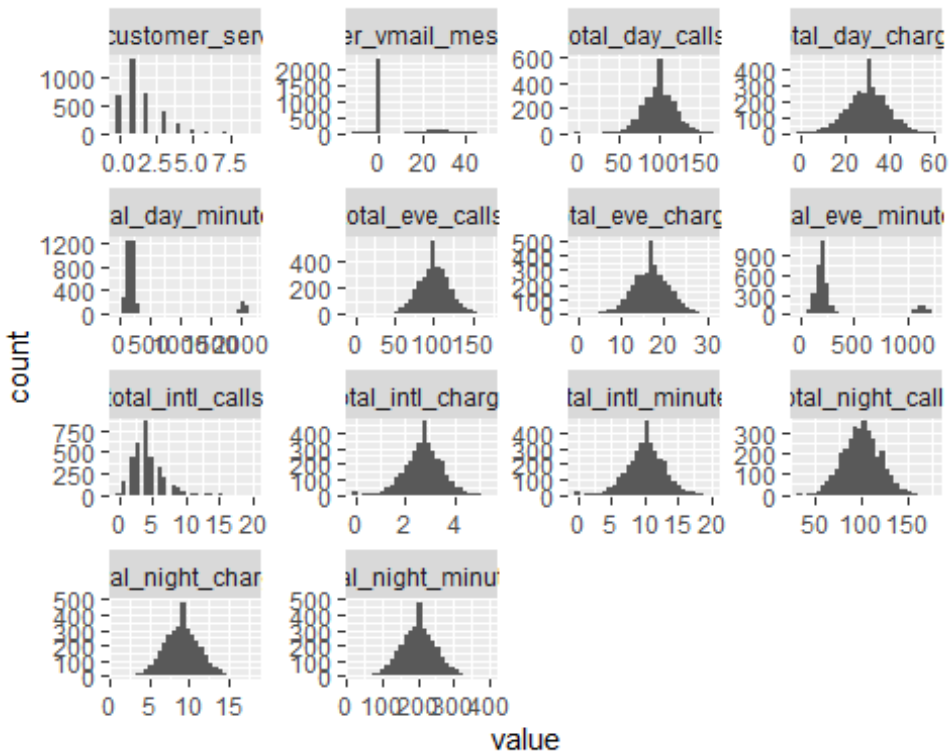
Exploratory Analysis

Perform Exploratory Analysis on Numerical Variables

```
# Explore the distribution of each variable using histograms
```


```
Churn_Train_Data %>%
  select_if(is.numeric) %>%
  gather() %>%
  ggplot(aes(value)) +
  facet_wrap(~key, scales = "free") +
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
#View(Churn_Train)
```

```
# Explore the relationship between pairs of variables using scatter plots
pairs( Churn_Train_Data %>%
  select_if(is.numeric))
```

Data Partition

Partitioning the Churn train data to training set of 75% and validation set of 25%.

```
set.seed(123)
```

```
Data_Partition<- createDataPartition(Churn_Train_Data$churn,p=0.75,list = FALSE)
```

```
train_data<- Churn_Train_Data[Data_Partition,]
```

```
Validation_Data<- Churn_Train_Data[-Data_Partition,]
```

Function For Threshold Identification

Created a function to find the threshold cutoff that balances Sensitivity, Specificity and Accuracy

```
FindThreshold <- function(actual, predict) {
```

```
  # Create a sequence of decimal numbers from 0 to 1 with a step of 0.1
```

```
  seq_decimal <- seq(from = 0.02, to = 1, by = 0.02)
```

```
  df <- data.frame(Threshold=0,Sensitivity=0,Specificity=0,Accuracy=0)
```

```
  rowNumber =1
```

```
  # Iterate over the sequence using a for loop
```

```
  for (i in seq_decimal) {
```

```
    predict_Lreg1<- ifelse(predict > i, 1, 0)
```

```
    x <- table(actual,predict_Lreg1)
```

```

df[rowNumber,1]=i
df[rowNumber,2]=(x[4]/(x[2]+x[4]))
df[rowNumber,3]=x[1]/(x[1]+x[3])
df[rowNumber,4]=(x[1]+x[4])/(x[2]+x[4]+x[1]+x[3])

rowNumber=rowNumber+1
}
return(df)
}

```

Create Function to calculate Metrics for the table output

```

MetricCalculation <- function(x){

  metriclist <- list(r1=c(
    TrueNegative=as.integer(x[1]),
    TruePositive=as.integer(x[4]),
    FalsePositive=as.integer(x[3]),
    FalseNegative=as.integer(x[2])),
    r2=c(
      sensitivityVal=as.double((x[4]/(x[2]+x[4])),4),
      SpecificityVal=round(x[1]/(x[1]+x[3]),4),
      Accuracy=(x[1]+x[4])/(x[2]+x[4]+x[1]+x[3])))

  return(metriclist)
}

```

Logistic Regression

```

Logistic_Model<- glm(churn~.,data = train_data, family = "binomial")
summary(Logistic_Model)

```

```

##
## Call:
## glm(formula = churn ~ ., family = "binomial", data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9657  -0.5058  -0.3404  -0.1929   3.0720
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -7.617943   0.831136  -9.166  < 2e-16 ***
## international_plan1  2.127728   0.164311  12.949  < 2e-16 ***
## voice_mail_plan1    -0.623079   0.341719  -1.823   0.06825 .
## number_vmail_messages -0.012730   0.011631  -1.095   0.27372
## total_day_minutes   -0.005818   0.002195  -2.650   0.00805 **
## total_day_calls     -0.001189   0.003307  -0.360   0.71919
## total_day_charge     0.100688   0.014257   7.063  1.64e-12 ***
## total_eve_minutes    0.011087   0.004314   2.570   0.01016 *
## total_eve_calls     -0.003379   0.003332  -1.014   0.31043
## total_eve_charge    -0.025673   0.050623  -0.507   0.61206

```

```
## total_night_minutes      0.049366    1.004785    0.049    0.96081
## total_night_calls        0.002408    0.003290    0.732    0.46427
## total_night_charge      -1.050252   22.327548   -0.047    0.96248
## total_intl_minutes       6.366234    6.333047    1.005    0.31478
## total_intl_calls        -0.123040    0.031510   -3.905   9.43e-05 ***
## total_intl_charge       -23.166968   23.453465   -0.988    0.32326
## number_customer_service_calls 0.512369    0.045670   11.219   < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 2071.8  on 2500  degrees of freedom
## Residual deviance: 1614.7  on 2484  degrees of freedom
## AIC: 1648.7
##
## Number of Fisher Scoring iterations: 6

predict_Lreg<- predict(Logistic_Model,Validation_Data, type = "response")
head(predict_Lreg)

##           1           2           7           15           17           18
## 0.02029239 0.36320969 0.09553940 0.07374417 0.04320038 0.02844423
```

Find the cutoff value

To identify the right threshold using the function built to calculate sensitivity, specificity and accuracy by changing the threshold value for identifying churn yes or no.

```
df <- FindThreshold(Validation_Data$churn,predict_Lreg)
df

##      Threshold Sensitivity Specificity Accuracy
## 1         0.02  0.94166667  0.1165730 0.2355769
## 2         0.04  0.90833333  0.2668539 0.3593750
## 3         0.06  0.87500000  0.4227528 0.4879808
## 4         0.08  0.83333333  0.5238764 0.5685096
## 5         0.10  0.80833333  0.6460674 0.6694712
## 6         0.12  0.76666667  0.6966292 0.7067308
## 7         0.14  0.71666667  0.7598315 0.7536058
## 8         0.16  0.66666667  0.7963483 0.7776442
## 9         0.18  0.63333333  0.8103933 0.7848558
## 10        0.20  0.56666667  0.8441011 0.8040865
## 11        0.22  0.52500000  0.8679775 0.8185096
## 12        0.24  0.49166667  0.8848315 0.8281250
## 13        0.26  0.45833333  0.9016854 0.8377404
## 14        0.28  0.41666667  0.9143258 0.8425481
## 15        0.30  0.40000000  0.9227528 0.8473558
## 16        0.32  0.38333333  0.9283708 0.8497596
## 17        0.34  0.35833333  0.9382022 0.8545673
## 18        0.36  0.30833333  0.9424157 0.8509615
```

## 19	0.38	0.26666667	0.9480337	0.8497596
## 20	0.40	0.26666667	0.9536517	0.8545673
## 21	0.42	0.24166667	0.9564607	0.8533654
## 22	0.44	0.22500000	0.9648876	0.8581731
## 23	0.46	0.21666667	0.9691011	0.8605769
## 24	0.48	0.20000000	0.9747191	0.8629808
## 25	0.50	0.16666667	0.9789326	0.8617788
## 26	0.52	0.15000000	0.9803371	0.8605769
## 27	0.54	0.14166667	0.9817416	0.8605769
## 28	0.56	0.14166667	0.9845506	0.8629808
## 29	0.58	0.12500000	0.9873596	0.8629808
## 30	0.60	0.11666667	0.9887640	0.8629808
## 31	0.62	0.11666667	0.9915730	0.8653846
## 32	0.64	0.11666667	0.9915730	0.8653846
## 33	0.66	0.10833333	0.9929775	0.8653846
## 34	0.68	0.10833333	0.9929775	0.8653846
## 35	0.70	0.09166667	0.9943820	0.8641827
## 36	0.72	0.07500000	0.9943820	0.8617788
## 37	0.74	0.06666667	0.9957865	0.8617788
## 38	0.76	0.05833333	0.9971910	0.8617788
## 39	0.78	0.05000000	0.9985955	0.8617788
## 40	0.80	0.04166667	0.9985955	0.8605769
## 41	0.82	0.03333333	0.9985955	0.8593750
## 42	0.84	0.02500000	0.9985955	0.8581731
## 43	0.86	0.01666667	0.9985955	0.8569712
## 44	0.88	0.01666667	0.9985955	0.8569712
## 45	0.90	0.00000000	0.9985955	0.8545673
## 46	0.92	NA	NA	NA
## 47	0.94	NA	NA	NA
## 48	0.96	NA	NA	NA
## 49	0.98	NA	NA	NA
## 50	1.00	NA	NA	NA

For Calculations we changed churn =Yes as 1 and churn=No as 0. Threshold 0.14 is selected after reviewing the threshold cut off table and metrics. Even though accuracy and Specificity are better with 0.16, we chose 0.14 as the customer churn(churn=Yes) costs telecom company more than the customers that would continue with the company(churn=No).

```
#print
df %>% filter(Threshold=="0.14")

##   Threshold Sensitivity Specificity Accuracy
## 1      0.14    0.7166667    0.7598315 0.7536058
```

we can pick the best threshold value to balance the prediction.

```
lg_threshold <- 0.14

predict_Lreg1<- ifelse(predict_Lreg > lg_threshold, 1, 0)
```

```

#predict_Lreg
tbl_Lreg <- table(Validation_Data$churn, predict_Lreg1)

tbl_Lreg

##      predict_Lreg1
##           0      1
##    0 541 171
##    1   34   86

#finding the accuracy
missing_class<- mean(predict_Lreg1 != Validation_Data$churn)
print(paste('Accuracy=', 1 - missing_class))

## [1] "Accuracy= 0.753605769230769"

```

ROC-AUC

```

ROC_Predict<- prediction(predict_Lreg1, Validation_Data$churn)
Roc_perform<- performance(ROC_Predict, measure = "tpr", x.measure = "fpr")

AUC_perform<- performance(ROC_Predict, measure = "auc")
AUC_perform<- AUC_perform@y.values[[1]]
AUC_perform

## [1] 0.7382491

```

Plot ROC Curve


```

plot.roc(Validation_Data$churn,predict_Lreg1)

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

```



```
Acc_perform<- performance(ROC_Predict, measure = "acc")
Acc_perform@y.values[[1]]
## [1] 0.8557692 0.7536058 0.1442308
```

Note: confusionMatrix function has provided sensitivity and specificity results in the reverse order

```
#printing the confusion matrix to see the prediction performance of the model
confusionMatrix(as.factor(predict_Lreg1),as.factor(Validation_Data$churn))

## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 541   34
##           1 171   86
##
##               Accuracy : 0.7536
##               95% CI   : (0.7229, 0.7825)
##       No Information Rate : 0.8558
##       P-Value [Acc > NIR] : 1
##
##               Kappa   : 0.3231
##
##  Mcnemar's Test P-Value : <2e-16
##
```

```
##          Sensitivity : 0.7598
##          Specificity : 0.7167
##          Pos Pred Value : 0.9409
##          Neg Pred Value : 0.3346
##          Prevalence : 0.8558
##          Detection Rate : 0.6502
##          Detection Prevalence : 0.6911
##          Balanced Accuracy : 0.7382
##
##          'Positive' Class : 0
##
```

```
CrossTable( Validation_Data$churn,predict_Lreg1,prop.chisq = F)
```

```
##
##
##      Cell Contents
## |-----|
## |                N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  832
##
```

```
##
##      Validation_Data$churn | predict_Lreg1
##      Validation_Data$churn |      0      1 | Row Total |
## -----|-----|-----|-----|
##              0 |      541      171 |      712 |
##              |      0.760      0.240 |      0.856 |
##              |      0.941      0.665 |      |
##              |      0.650      0.206 |      |
## -----|-----|-----|-----|
##              1 |       34       86 |      120 |
##              |      0.283      0.717 |      0.144 |
##              |      0.059      0.335 |      |
##              |      0.041      0.103 |      |
## -----|-----|-----|-----|
##      Column Total |      575      257 |      832 |
##              |      0.691      0.309 |      |
## -----|-----|-----|-----|
##
##
```

Snapshot of Final Metrics for the validation data using Logistic regression

```
MetricCalculation(tbl_Lreg)
```

```
## $r1
## TrueNegative TruePositive FalsePositive FalseNegative
##          541          86          171          34
##
## $r2
## sensitivityVal SpecificityVal      Accuracy
##      0.7166667      0.7598000      0.7536058
```

KNN model

```
#set.seed(125)
```

```
KNN_Model<- knn(train = train_data[,1:17],test =Validation_Data[,1:17], cl=
train_data$churn,
                k=60 ,prob = TRUE )
```

```
probs <- attr(KNN_Model,"prob")
```

```
dfWithProbYes <- data.frame(KNN_Model,probs)
```

```
dfWithProbYes <- ifelse(dfWithProbYes$KNN_Model==1,dfWithProbYes$probs,1-
dfWithProbYes$probs)
```

```
head(dfWithProbYes)
```

```
## [1] 0.20000000 0.58333333 0.09523810 0.10000000 0.14516129 0.08196721
```

Find the cutoff value

```
df_knn <- FindThreshold(Validation_Data$churn,dfWithProbYes)
df_knn
```

```
##      Threshold Sensitivity Specificity Accuracy
## 1          0.02 1.000000000 0.01264045 0.1550481
## 2          0.04 0.983333333 0.04073034 0.1766827
## 3          0.06 0.958333333 0.09269663 0.2175481
## 4          0.08 0.866666667 0.19943820 0.2956731
## 5          0.10 0.750000000 0.46207865 0.5036058
## 6          0.12 0.658333333 0.61095506 0.6177885
## 7          0.14 0.608333333 0.70365169 0.6899038
## 8          0.16 0.541666667 0.79213483 0.7560096
## 9          0.18 0.475000000 0.87500000 0.8173077
## 10         0.20 0.425000000 0.93960674 0.8653846
## 11         0.22 0.416666667 0.95224719 0.8750000
## 12         0.24 0.416666667 0.96207865 0.8834135
## 13         0.26 0.408333333 0.97752809 0.8954327
## 14         0.28 0.391666667 0.97893258 0.8942308
## 15         0.30 0.358333333 0.98595506 0.8954327
## 16         0.32 0.333333333 0.98735955 0.8930288
## 17         0.34 0.300000000 0.99157303 0.8918269
## 18         0.36 0.283333333 0.99297753 0.8906250
```



```
## 19      0.38 0.258333333 0.99438202 0.8882212
## 20      0.40 0.166666667 1.00000000 0.8798077
## 21      0.42 0.158333333 1.00000000 0.8786058
## 22      0.44 0.150000000 1.00000000 0.8774038
## 23      0.46 0.133333333 1.00000000 0.8750000
## 24      0.48 0.125000000 1.00000000 0.8737981
## 25      0.50 0.116666667 1.00000000 0.8725962
## 26      0.52 0.100000000 1.00000000 0.8701923
## 27      0.54 0.083333333 1.00000000 0.8677885
## 28      0.56 0.083333333 1.00000000 0.8677885
## 29      0.58 0.066666667 1.00000000 0.8653846
## 30      0.60 0.050000000 1.00000000 0.8629808
## 31      0.62 0.033333333 1.00000000 0.8605769
## 32      0.64 0.025000000 1.00000000 0.8593750
## 33      0.66 0.025000000 1.00000000 0.8593750
## 34      0.68 0.025000000 1.00000000 0.8593750
## 35      0.70 0.016666667 1.00000000 0.8581731
## 36      0.72 0.008333333 1.00000000 0.8569712
## 37      0.74      NA      NA      NA
## 38      0.76      NA      NA      NA
## 39      0.78      NA      NA      NA
## 40      0.80      NA      NA      NA
## 41      0.82      NA      NA      NA
## 42      0.84      NA      NA      NA
## 43      0.86      NA      NA      NA
## 44      0.88      NA      NA      NA
## 45      0.90      NA      NA      NA
## 46      0.92      NA      NA      NA
## 47      0.94      NA      NA      NA
## 48      0.96      NA      NA      NA
## 49      0.98      NA      NA      NA
## 50      1.00      NA      NA      NA
```

Threshold selection from the combination found in the df_knn

```
#print
df_knn %>% filter(Threshold=="0.14")

##   Threshold Sensitivity Specificity Accuracy
## 1      0.14    0.6083333    0.7036517 0.6899038
```

we can pick the best threshold value to balance the prediction using the combination of sensitivity, specificity and accuracy.

```
knn_threshold <- 0.14

predict_knn<- ifelse(dfWithProbYes > knn_threshold, 1, 0)
#predict_Lreg
tbl_knn <- table(Validation_Data$churn, predict_knn)

tbl_knn
```

```
##      predict_knn
##      0      1
##      0 501 211
##      1  47  73
```

Note: confusionMatrix function has provided sensitivity and specificity results in the reverse order

```
confusionMatrix(as.factor(predict_knn),Validation_Data$churn)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 501  47
##              1 211  73
##
##              Accuracy : 0.6899
##              95% CI : (0.6572, 0.7212)
##      No Information Rate : 0.8558
##      P-Value [Acc > NIR] : 1
##
##              Kappa : 0.1989
##
## Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.7037
##              Specificity : 0.6083
##              Pos Pred Value : 0.9142
##              Neg Pred Value : 0.2570
##              Prevalence : 0.8558
##              Detection Rate : 0.6022
##      Detection Prevalence : 0.6587
##              Balanced Accuracy : 0.6560
##
##              'Positive' Class : 0
##
```

```
CrossTable(Validation_Data$churn,predict_knn)
```

```
##
##
##      Cell Contents
## |-----|
## |              N |
## | Chi-square contribution |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
```

```
##
## Total Observations in Table: 832
##
##
## predict_knn
## Validation_Data$churn |      0      1 | Row Total |
## -----|-----|-----|
##           0 |      501      211 |      712 |
##           |      2.189      4.223 |      |
##           |      0.704      0.296 |      0.856 |
##           |      0.914      0.743 |      |
##           |      0.602      0.254 |      |
## -----|-----|-----|
##           1 |       47       73 |      120 |
##           |     12.987     25.059 |      |
##           |      0.392      0.608 |      0.144 |
##           |      0.086      0.257 |      |
##           |      0.056      0.088 |      |
## -----|-----|-----|
##      Column Total |      548      284 |      832 |
##           |      0.659      0.341 |      |
## -----|-----|-----|
##
##
```

```
MetricCalculation(tbl_knn)
```

```
## $r1
## TrueNegative TruePositive FalsePositive FalseNegative
##           501           73           211           47
##
## $r2
## sensitivityVal SpecificityVal      Accuracy
##           0.6083333           0.7037000           0.6899038
```

Decision Tree Model

Build and Prune Decision tree using the best CP value

```
Dec_tree_model<- rpart(churn~., data = train_data, method = "class",
                        control = rpart.control(minsplit = 40))
```

#Find the best cp value for pruning decision tree

```
Best_cp<-
```

```
Dec_tree_model$cptable[which.min(Dec_tree_model$cptable[, "xerror"]), "CP"]
```

#Prune decision tree based on the cp value identified by building base model

```
Dec_tree_model <- prune(Dec_tree_model, cp=Best_cp)
```

```
summary(Dec_tree_model)
```

```

## Call:
## rpart(formula = churn ~ ., data = train_data, method = "class",
##       control = rpart.control(minsplit = 40))
##       n= 2501
##
##              CP nsplit rel error      xerror      xstd
## 1 0.07713499      0 1.0000000 1.0000000 0.04852815
## 2 0.05647383      4 0.6914601 0.8099174 0.04437224
## 3 0.03030303      7 0.4931129 0.5399449 0.03702542
## 4 0.01652893      8 0.4628099 0.5261708 0.03658973
## 5 0.01101928     10 0.4297521 0.5179063 0.03632480
## 6 0.01000000     11 0.4187328 0.5151515 0.03623590
##
## Variable importance
##              total_day_charge number_customer_service_calls
##                      21                                12
##      international_plan              total_intl_charge
##                      9                                9
##      total_intl_minutes              total_day_minutes
##                      9                                8
##      total_intl_calls              total_eve_charge
##                      8                                8
##      total_eve_minutes              voice_mail_plan
##                      5                                5
##      number_vmail_messages              total_night_charge
##                      5                                1
##      total_night_minutes              total_day_calls
##                      1                                1
##
## Node number 1: 2501 observations,      complexity param=0.07713499
## predicted class=0 expected loss=0.1451419 P(node) =1
## class counts: 2138 363
## probabilities: 0.855 0.145
## left son=2 (2355 obs) right son=3 (146 obs)
## Primary splits:
##      total_day_charge < 44.985 to the left,
improve=59.233520, (0 missing)
##      number_customer_service_calls < 3.5 to the left,
improve=56.684250, (0 missing)
##      international_plan splits as LR,
improve=48.042920, (0 missing)
##      total_day_minutes < 223.25 to the left,
improve=16.822220, (0 missing)
##      voice_mail_plan splits as RL, improve=
7.072234, (0 missing)
##
## Node number 2: 2355 observations,      complexity param=0.07713499
## predicted class=0 expected loss=0.1180467 P(node) =0.9416234
## class counts: 2077 278
## probabilities: 0.882 0.118

```

```

## left son=4 (2168 obs) right son=5 (187 obs)
## Primary splits:
## number_customer_service_calls < 3.5 to the left,
improve=60.101180, (0 missing)
## international_plan splits as LR,
improve=46.007340, (0 missing)
## total_day_charge < 37.95 to the left, improve=
7.622099, (0 missing)
## total_intl_calls < 2.5 to the right, improve=
6.770627, (0 missing)
## total_intl_minutes < 13.15 to the left, improve=
6.020005, (0 missing)
## Surrogate splits:
## total_day_calls < 38 to the right, agree=0.921, adj=0.005, (0
split)
##
## Node number 3: 146 observations, complexity param=0.07713499
## predicted class=1 expected loss=0.4178082 P(node) =0.05837665
## class counts: 61 85
## probabilities: 0.418 0.582
## left son=6 (38 obs) right son=7 (108 obs)
## Primary splits:
## voice_mail_plan splits as RL, improve=23.369500, (0
missing)
## number_vmail_messages < 6.5 to the right, improve=22.277580, (0
missing)
## total_eve_minutes < 185.45 to the left, improve=12.985400, (0
missing)
## total_eve_charge < 17.075 to the left, improve=11.506940, (0
missing)
## total_day_charge < 53.84 to the left, improve= 5.836558, (0
missing)
## Surrogate splits:
## number_vmail_messages < 6.5 to the right, agree=0.993,
adj=0.974, (0 split)
## total_eve_minutes < 1126.15 to the right, agree=0.767,
adj=0.105, (0 split)
## total_day_minutes < 2080.3 to the right, agree=0.760,
adj=0.079, (0 split)
## total_night_minutes < 119.55 to the left, agree=0.760,
adj=0.079, (0 split)
## total_night_charge < 5.38 to the left, agree=0.760,
adj=0.079, (0 split)
##
## Node number 4: 2168 observations, complexity param=0.05647383
## predicted class=0 expected loss=0.08487085 P(node) =0.8668533
## class counts: 1984 184
## probabilities: 0.915 0.085
## left son=8 (1952 obs) right son=9 (216 obs)
## Primary splits:

```

```

##      international_plan splits as  LR,          improve=45.707680, (0
missing)
##      total_day_charge  < 37.95  to the left,  improve=12.804600, (0
missing)
##      total_eve_charge  < 21.175  to the left,  improve= 6.107379, (0
missing)
##      total_intl_calls  < 2.5     to the right, improve= 4.941350, (0
missing)
##      total_intl_minutes < 13.15  to the left,  improve= 4.925431, (0
missing)
##      Surrogate splits:
##      total_eve_minutes < 1234.3  to the left,  agree=0.901, adj=0.009, (0
split)
##
## Node number 5: 187 observations,    complexity param=0.07713499
##      predicted class=1  expected loss=0.4973262  P(node) =0.07477009
##      class counts:    93    94
##      probabilities: 0.497 0.503
##      left son=10 (109 obs) right son=11 (78 obs)
##      Primary splits:
##      total_day_charge          < 27.685  to the right,
improve=36.465100, (0 missing)
##      total_day_minutes          < 162.85  to the right,
improve=32.460370, (0 missing)
##      total_eve_minutes          < 264.65  to the right, improve=
7.583369, (0 missing)
##      total_eve_charge          < 15.825  to the right, improve=
6.881493, (0 missing)
##      number_customer_service_calls < 4.5    to the left,  improve=
4.344669, (0 missing)
##      Surrogate splits:
##      total_day_minutes  < 162.85  to the right, agree=0.968, adj=0.923,
(0 split)
##      total_intl_calls   < 2.5     to the right, agree=0.604, adj=0.051,
(0 split)
##      total_eve_calls    < 86.5    to the right, agree=0.599, adj=0.038,
(0 split)
##      total_night_minutes < 91.15  to the right, agree=0.599, adj=0.038,
(0 split)
##      total_night_charge  < 4.1     to the right, agree=0.599, adj=0.038,
(0 split)
##
## Node number 6: 38 observations
##      predicted class=0  expected loss=0.1052632  P(node) =0.01519392
##      class counts:    34     4
##      probabilities: 0.895 0.105
##
## Node number 7: 108 observations,    complexity param=0.03030303
##      predicted class=1  expected loss=0.25  P(node) =0.04318273
##      class counts:    27    81

```

```

##   probabilities: 0.250 0.750
##   left son=14 (33 obs) right son=15 (75 obs)
##   Primary splits:
##       total_eve_minutes < 183.75 to the left, improve=16.500000, (0
missing)
##       total_eve_charge < 15.695 to the left, improve=15.564710, (0
missing)
##       total_night_minutes < 169.6 to the left, improve= 3.976364, (0
missing)
##       total_night_charge < 7.63 to the left, improve= 3.976364, (0
missing)
##       total_day_charge < 47.34 to the left, improve= 3.146897, (0
missing)
##   Surrogate splits:
##       total_eve_charge < 15.545 to the left, agree=0.944, adj=0.818,
(0 split)
##       total_intl_minutes < 3.35 to the left, agree=0.713, adj=0.061,
(0 split)
##       total_intl_charge < 0.905 to the left, agree=0.713, adj=0.061,
(0 split)
##       total_eve_calls < 117.5 to the right, agree=0.704, adj=0.030,
(0 split)
##       total_intl_calls < 1.5 to the left, agree=0.704, adj=0.030,
(0 split)
##
## Node number 8: 1952 observations, complexity param=0.01652893
## predicted class=0 expected loss=0.05071721 P(node) =0.7804878
## class counts: 1853 99
## probabilities: 0.949 0.051
## left son=16 (1699 obs) right son=17 (253 obs)
## Primary splits:
##       total_day_charge < 37.95 to the left, improve=9.398504, (0
missing)
##       total_eve_charge < 20.855 to the left, improve=4.728740, (0
missing)
##       total_day_minutes < 223.25 to the left, improve=3.461978, (0
missing)
##       total_eve_minutes < 208.85 to the left, improve=2.903298, (0
missing)
##       total_night_minutes < 191.3 to the left, improve=1.369124, (0
missing)
##   Surrogate splits:
##       total_day_minutes < 223.25 to the left, agree=0.894, adj=0.182, (0
split)
##
## Node number 9: 216 observations, complexity param=0.05647383
## predicted class=0 expected loss=0.3935185 P(node) =0.08636545
## class counts: 131 85
## probabilities: 0.606 0.394
## left son=18 (175 obs) right son=19 (41 obs)

```

```

## Primary splits:
## total_intl_calls < 2.5 to the right, improve=37.227570, (0
missing)
## total_intl_minutes < 13.05 to the left, improve=30.726160, (0
missing)
## total_intl_charge < 3.52 to the left, improve=30.726160, (0
missing)
## total_eve_calls < 107.5 to the right, improve= 3.296622, (0
missing)
## total_night_calls < 74.5 to the left, improve= 3.132155, (0
missing)
## Surrogate splits:
## total_day_calls < 49 to the right, agree=0.819, adj=0.049, (0
split)
##
## Node number 10: 109 observations, complexity param=0.01101928
## predicted class=0 expected loss=0.2385321 P(node) =0.04358257
## class counts: 83 26
## probabilities: 0.761 0.239
## left son=20 (95 obs) right son=21 (14 obs)
## Primary splits:
## total_eve_charge < 12 to the right, improve=5.251969, (0
missing)
## total_eve_minutes < 144.55 to the right, improve=4.192484, (0
missing)
## total_day_charge < 29.88 to the right, improve=3.350628, (0
missing)
## total_day_minutes < 195.2 to the right, improve=2.353577, (0
missing)
## total_night_calls < 116.5 to the left, improve=1.878798, (0
missing)
## Surrogate splits:
## total_eve_minutes < 141.15 to the right, agree=0.972, adj=0.786, (0
split)
##
## Node number 11: 78 observations
## predicted class=1 expected loss=0.1282051 P(node) =0.03118752
## class counts: 10 68
## probabilities: 0.128 0.872
##
## Node number 14: 33 observations
## predicted class=0 expected loss=0.3333333 P(node) =0.01319472
## class counts: 22 11
## probabilities: 0.667 0.333
##
## Node number 15: 75 observations
## predicted class=1 expected loss=0.06666667 P(node) =0.029988
## class counts: 5 70
## probabilities: 0.067 0.933
##

```



```

## Node number 16: 1699 observations
##   predicted class=0   expected loss=0.0317834   P(node) =0.6793283
##   class counts:  1645    54
##   probabilities: 0.968 0.032
##
## Node number 17: 253 observations,   complexity param=0.01652893
##   predicted class=0   expected loss=0.1778656   P(node) =0.1011595
##   class counts:    208    45
##   probabilities: 0.822 0.178
##   left son=34 (217 obs) right son=35 (36 obs)
##   Primary splits:
##       total_eve_charge      < 22.08   to the left,   improve=20.056610, (0
missing)
##       total_eve_minutes    < 240.15  to the left,   improve=12.632570, (0
missing)
##       voice_mail_plan      splits as  RL,          improve= 3.785385, (0
missing)
##       number_vmail_messages < 5.5     to the right, improve= 3.582237, (0
missing)
##       total_night_minutes  < 181.15  to the left,   improve= 3.015811, (0
missing)
##   Surrogate splits:
##       total_eve_minutes    < 259.8    to the left,   agree=0.877, adj=0.139,
(0 split)
##       total_day_calls      < 135.5    to the left,   agree=0.862, adj=0.028,
(0 split)
##       total_intl_minutes   < 2.05     to the right,  agree=0.862, adj=0.028,
(0 split)
##       total_intl_charge    < 0.555    to the right,  agree=0.862, adj=0.028,
(0 split)
##
## Node number 18: 175 observations,   complexity param=0.05647383
##   predicted class=0   expected loss=0.2514286   P(node) =0.06997201
##   class counts:    131    44
##   probabilities: 0.749 0.251
##   left son=36 (144 obs) right son=37 (31 obs)
##   Primary splits:
##       total_intl_minutes   < 13.05    to the left,   improve=42.221510, (0
missing)
##       total_intl_charge    < 3.52     to the left,   improve=42.221510, (0
missing)
##       total_eve_charge     < 14.11    to the left,   improve= 2.860896, (0
missing)
##       total_night_minutes  < 216.7    to the right,  improve= 2.472771, (0
missing)
##       total_night_charge   < 9.75     to the right,  improve= 2.472771, (0
missing)
##   Surrogate splits:
##       total_intl_charge    < 3.52     to the left,   agree=1.000, adj=1.000, (0
split)

```

```

##      total_day_minutes < 55.3      to the right, agree=0.834, adj=0.065, (0
split)
##      total_day_charge < 8.92      to the right, agree=0.829, adj=0.032, (0
split)
##
## Node number 19: 41 observations
##   predicted class=1   expected loss=0   P(node) =0.01639344
##   class counts:      0    41
##   probabilities: 0.000 1.000
##
## Node number 20: 95 observations
##   predicted class=0   expected loss=0.1789474   P(node) =0.03798481
##   class counts:      78    17
##   probabilities: 0.821 0.179
##
## Node number 21: 14 observations
##   predicted class=1   expected loss=0.3571429   P(node) =0.005597761
##   class counts:       5     9
##   probabilities: 0.357 0.643
##
## Node number 34: 217 observations
##   predicted class=0   expected loss=0.09677419   P(node) =0.08676529
##   class counts:     196    21
##   probabilities: 0.903 0.097
##
## Node number 35: 36 observations
##   predicted class=1   expected loss=0.3333333   P(node) =0.01439424
##   class counts:      12    24
##   probabilities: 0.333 0.667
##
## Node number 36: 144 observations
##   predicted class=0   expected loss=0.09027778   P(node) =0.05757697
##   class counts:     131    13
##   probabilities: 0.910 0.090
##
## Node number 37: 31 observations
##   predicted class=1   expected loss=0   P(node) =0.01239504
##   class counts:       0    31
##   probabilities: 0.000 1.000

```

#Probability prediction

```

Probability_DecisionTree <- predict(Dec_tree_model, newdata =
Validation_Data, type = "prob")

```

#calculating AUC Value

```

ROC_Predict_dt<- prediction(Probability_DecisionTree[,2],
Validation_Data$churn)
Roc_perform_dt<- performance(ROC_Predict_dt, measure = "tpr", x.measure =
"fpr")

```

```

AUC_perform_dt<- performance(ROC_Predict_dt, measure = "auc")
AUC_perform_dt<- AUC_perform_dt@y.values[[1]]
AUC_perform_dt

## [1] 0.875357

df_dt <- FindThreshold(Validation_Data$churn,Probability_DecisionTree[,2] )
df_dt

##      Threshold Sensitivity Specificity Accuracy
## 1          0.02           NA           NA         NA
## 2          0.04    0.8416667    0.7640449 0.7752404
## 3          0.06    0.8416667    0.7640449 0.7752404
## 4          0.08    0.8416667    0.7640449 0.7752404
## 5          0.10    0.7166667    0.9283708 0.8978365
## 6          0.12    0.7000000    0.9452247 0.9098558
## 7          0.14    0.7000000    0.9452247 0.9098558
## 8          0.16    0.7000000    0.9452247 0.9098558
## 9          0.18    0.6666667    0.9719101 0.9278846
## 10         0.20    0.6666667    0.9719101 0.9278846
## 11         0.22    0.6666667    0.9719101 0.9278846
## 12         0.24    0.6666667    0.9719101 0.9278846
## 13         0.26    0.6666667    0.9719101 0.9278846
## 14         0.28    0.6666667    0.9719101 0.9278846
## 15         0.30    0.6666667    0.9719101 0.9278846
## 16         0.32    0.6666667    0.9719101 0.9278846
## 17         0.34    0.6166667    0.9775281 0.9254808
## 18         0.36    0.6166667    0.9775281 0.9254808
## 19         0.38    0.6166667    0.9775281 0.9254808
## 20         0.40    0.6166667    0.9775281 0.9254808
## 21         0.42    0.6166667    0.9775281 0.9254808
## 22         0.44    0.6166667    0.9775281 0.9254808
## 23         0.46    0.6166667    0.9775281 0.9254808
## 24         0.48    0.6166667    0.9775281 0.9254808
## 25         0.50    0.6166667    0.9775281 0.9254808
## 26         0.52    0.6166667    0.9775281 0.9254808
## 27         0.54    0.6166667    0.9775281 0.9254808
## 28         0.56    0.6166667    0.9775281 0.9254808
## 29         0.58    0.6166667    0.9775281 0.9254808
## 30         0.60    0.6166667    0.9775281 0.9254808
## 31         0.62    0.6166667    0.9775281 0.9254808
## 32         0.64    0.6166667    0.9775281 0.9254808
## 33         0.66    0.5916667    0.9775281 0.9218750
## 34         0.68    0.5333333    0.9845506 0.9194712
## 35         0.70    0.5333333    0.9845506 0.9194712
## 36         0.72    0.5333333    0.9845506 0.9194712
## 37         0.74    0.5333333    0.9845506 0.9194712
## 38         0.76    0.5333333    0.9845506 0.9194712
## 39         0.78    0.5333333    0.9845506 0.9194712
## 40         0.80    0.5333333    0.9845506 0.9194712

```

```
## 41      0.82    0.5333333    0.9845506 0.9194712
## 42      0.84    0.5333333    0.9845506 0.9194712
## 43      0.86    0.5333333    0.9845506 0.9194712
## 44      0.88    0.3750000    0.9929775 0.9038462
## 45      0.90    0.3750000    0.9929775 0.9038462
## 46      0.92    0.3750000    0.9929775 0.9038462
## 47      0.94    0.1333333    1.0000000 0.8750000
## 48      0.96    0.1333333    1.0000000 0.8750000
## 49      0.98    0.1333333    1.0000000 0.8750000
## 50      1.00           NA           NA           NA
```

we can pick the best threshold value to balance the prediction.

So, 0.10 is selected which has better sensitivity and specificity combined resulted in better accuracy as well.

```
dt_threshold <- 0.10
predict_dt<- ifelse(Probability_DecisionTree[,2] > dt_threshold, 1, 0)

tbl_dt <- table(Validation_Data$churn, predict_dt)

tbl_dt

##      predict_dt
##           0    1
##    0 661   51
##    1   34   86
```

Note: confusionMatrix function has provided sensitivity and specificity results in the reverse order

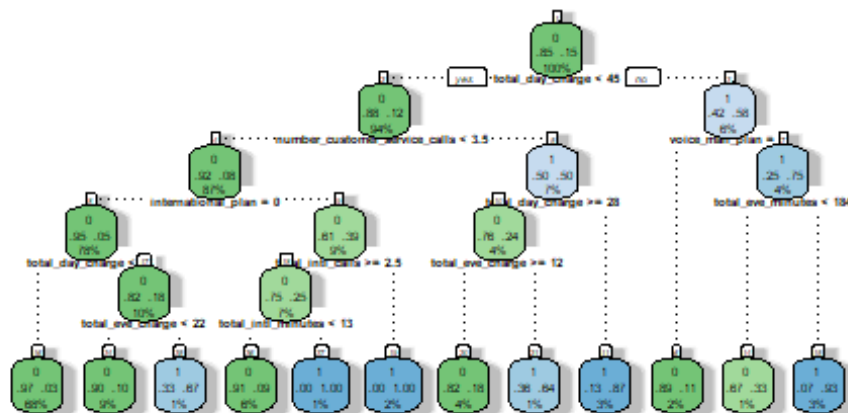
#printing the confusion matrix to see the prediction performance of the model
 confusionMatrix(as.factor(predict_dt),as.factor(Validation_Data\$churn))

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##           0 661   34
##           1  51   86
##
##              Accuracy : 0.8978
##              95% CI   : (0.8752, 0.9176)
##      No Information Rate : 0.8558
##      P-Value [Acc > NIR] : 0.0001938
##
##              Kappa   : 0.6092
##
##  Mcnemar's Test P-Value : 0.0826623
##
##              Sensitivity : 0.9284
```

```
##          Specificity : 0.7167
##          Pos Pred Value : 0.9511
##          Neg Pred Value : 0.6277
##          Prevalence : 0.8558
##          Detection Rate : 0.7945
##          Detection Prevalence : 0.8353
##          Balanced Accuracy : 0.8225
##
##          'Positive' Class : 0
##
```

Plot Decision Tree Model

```
fancyRpartPlot(Dec_tree_model,cex=0.4)
```



Rattle 2023-Apr-30 22:28:20 peddi

```
CrossTable(Validation_Data$churn,predict_dt,prop.chisq = F)
```

```
##
##
##      Cell Contents
## |-----|
## |               N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
```

```
## Total Observations in Table: 832
##
##
##      | predict_dt
## Validation_Data$churn |      0      1 | Row Total |
## -----|-----|-----|-----|
##           0 |      661      51 |      712 |
##           |      0.928      0.072 |      0.856 |
##           |      0.951      0.372 |
##           |      0.794      0.061 |
## -----|-----|-----|
##           1 |       34      86 |      120 |
##           |      0.283      0.717 |      0.144 |
##           |      0.049      0.628 |
##           |      0.041      0.103 |
## -----|-----|-----|
##      Column Total |      695      137 |      832 |
##           |      0.835      0.165 |
## -----|-----|-----|
##
##
```

```
MetricCalculation(tbl_dt)
```

```
## $r1
## TrueNegative TruePositive FalsePositive FalseNegative
##           661           86           51           34
##
## $r2
## sensitivityVal SpecificityVal      Accuracy
##      0.7166667      0.9284000      0.8978365
```

Metrics output for model selection

```
data.frame("Prediction Models"= c("Logistic Regression","KNN","Decision
Tree"))

,Sensitivity=c(MetricCalculation(tbl_Lreg)[[2]][1],MetricCalculation(tbl_knn)
[[2]][1],MetricCalculation(tbl_dt)[[2]][1])

,Specificity=c(MetricCalculation(tbl_Lreg)[[2]][2],MetricCalculation(tbl_knn)
[[2]][2],MetricCalculation(tbl_dt)[[2]][2])

,Accuracy=c(MetricCalculation(tbl_Lreg)[[2]][3],MetricCalculation(tbl_knn)[[2]
][3],MetricCalculation(tbl_dt)[[2]][3])
)

##      Prediction.Models Sensitivity Specificity Accuracy
## 1 Logistic Regression   0.7166667      0.7598 0.7536058
## 2              KNN     0.6083333      0.7037 0.6899038
## 3      Decision Tree    0.7166667      0.9284 0.8978365
```

Model Selection

After comparing the sensitivity, specificity, and accuracy metrics of the three models built in the project, the Decision tree and Logistic regression have the best sensitivity. Even though decision tree and logistic regression have similar sensitivity, the specificity and accuracy of logistic regression are lower, which could cost the telecom company significantly higher costs as the target customers for reducing churn will increase if specificity is low. So, the Decision tree has the best sensitivity, specificity, and accuracy, resulting in lower marketing costs to retain ABC telecom customers that are more likely to churn.

Decision Tree is the best model of three models built as part of the project.

Predict Churn for Customers_To_Predict

#finding the number of rows in the customers to predict dataset.

```
count(Customers_To_Predict)
```

```
## # A tibble: 1 × 1
```

```
##       n
```

```
##   <int>
```

```
## 1  1600
```

#printing the summary of the testset

```
summary(Customers_To_Predict)
```

```
##      state      account_length      area_code      international_plan
## Length:1600    Min.   : 1.00    Length:1600    Length:1600
## Class :character 1st Qu.: 71.00    Class :character Class :character
## Mode  :character Median : 98.00    Mode  :character Mode  :character
##                      Mean   : 98.52
##                      3rd Qu.:126.00
##                      Max.   :238.00
## voice_mail_plan  number_vmail_messages total_day_minutes
total_day_calls
## Length:1600    Min.   : 0.000    Min.   : 6.6    Min.   : 34.00
## Class :character 1st Qu.: 0.000    1st Qu.:143.8    1st Qu.: 86.00
## Mode  :character Median : 0.000    Median :180.9    Median : 99.00
##                      Mean   : 7.043    Mean   :181.6    Mean   : 99.06
##                      3rd Qu.: 0.000    3rd Qu.:215.9    3rd Qu.:112.00
##                      Max.   :52.000    Max.   :351.5    Max.   :160.00
## total_day_charge total_eve_minutes total_eve_calls total_eve_charge
## Min.   : 1.12    Min.   : 22.3    Min.   : 38.0    Min.   : 1.90
## 1st Qu.:24.45    1st Qu.:165.8    1st Qu.: 88.0    1st Qu.:14.10
## Median :30.76    Median :199.9    Median :101.0    Median :17.00
## Mean   :30.87    Mean   :199.6    Mean   :100.6    Mean   :16.96
## 3rd Qu.:36.70    3rd Qu.:231.8    3rd Qu.:114.0    3rd Qu.:19.70
## Max.   :59.76    Max.   :359.3    Max.   :169.0    Max.   :30.54
## total_night_minutes total_night_calls total_night_charge
total_intl_minutes
## Min.   : 0.0    Min.   : 0.00    Min.   : 0.000    Min.   : 0.00
```

```
## 1st Qu.:166.6      1st Qu.: 86.00      1st Qu.: 7.500      1st Qu.: 8.60
## Median :199.2      Median : 99.00      Median : 8.960      Median :10.40
## Mean   :199.2      Mean   : 99.45      Mean   : 8.963      Mean   :10.32
## 3rd Qu.:232.4      3rd Qu.:113.00     3rd Qu.:10.463     3rd Qu.:12.00
## Max.   :381.6      Max.   :170.00     Max.   :17.170     Max.   :19.70
## total_intl_calls total_intl_charge number_customer_service_calls
## Min.    : 0.000    Min.    :0.000    Min.    :0.000
## 1st Qu.: 3.000    1st Qu.:2.320    1st Qu.:1.000
## Median : 4.000    Median :2.810    Median :1.000
## Mean   : 4.356    Mean   :2.786    Mean   :1.583
## 3rd Qu.: 5.000    3rd Qu.:3.240    3rd Qu.:2.000
## Max.   :19.000    Max.   :5.320    Max.   :7.000
```

#checking for any null values

```
any(is.na(Customers_To_Predict))
```

```
## [1] FALSE
```

#finding the percentage of null values

```
colMeans(is.na(Customers_To_Predict))*100
```

```
##              state              account_length
##              0                      0
##              area_code            international_plan
##              0                      0
##              voice_mail_plan      number_vmail_messages
##              0                      0
##              total_day_minutes    total_day_calls
##              0                      0
##              total_day_charge      total_eve_minutes
##              0                      0
##              total_eve_calls      total_eve_charge
##              0                      0
##              total_night_minutes  total_night_calls
##              0                      0
##              total_night_charge    total_intl_minutes
##              0                      0
##              total_intl_calls      total_intl_charge
##              0                      0
## number_customer_service_calls
##              0
```

```
Customers_To_Predict$international_plan<-
```

```
ifelse(Customers_To_Predict$international_plan=="yes",1,0)
```

```
Customers_To_Predict$voice_mail_plan<-
```

```
ifelse(Customers_To_Predict$voice_mail_plan=="yes",1,0)
```

#converting the numbers into factors

```
Customers_To_Predict$international_plan<-
```

```
as.factor(Customers_To_Predict$international_plan)
```



```

Customers_To_Predict$voice_mail_plan<-
as.factor(Customers_To_Predict$voice_mail_plan)

#For Logistic regression change the below code with
predict(Logistic_Model,newdata = Customers_To_Predict, type = "response")

#running the predictions using the best decision model built.
predict_BestTree_Model<- predict(Dec_tree_model, newdata =
Customers_To_Predict, type = "prob" )

Using the threshold identified from the decision tree built on the train data.
Customer_to_predict output data churns are determined from the probabilities.
PredictedChurn <- ifelse(predict_BestTree_Model[,2] > dt_threshold, 1, 0)

summary(as.factor(PredictedChurn))

##      0      1
## 1329   271

#Save customer to predict data with predicted churn to the
decisiontreeoutput.csv
write.csv(cbind(Customers_To_Predict,PredictedChurn),file='DecisionTreeOutput
.csv')

```